

# Understanding LlamalIndex

---

- Understanding LlamalIndex

Ram N Sangwan

# LLM Fine Tuning?

---

LLMs are not trained on **your** data, which may be private or specific to the problem you're trying to solve.

It's behind APIs, in SQL databases, or trapped in PDFs and slide decks.

You may choose to **fine-tune** a LLM with your data, but:

- Training a LLM is **expensive**.
- Due to the cost to train, it's **hard to update** a LLM with latest information.
- **Observability** is lacking. When you ask a LLM a question, it's not obvious how the LLM arrived at its answer.

# Challenges with Naive RAG (Response Quality)

- **Bad Retrieval**


- **Low Precision:** Not all chunks in retrieved set are relevant
  - Hallucination + Lost in the Middle Problems
- **Low Recall:** Not all relevant chunks are retrieved.
  - Lacks enough context for LLM to synthesize an answer
- **Outdated information:** The data is redundant or out of date.

- **Bad Response Generation**

- **Hallucination:** Model makes up an answer that isn't in the context.
- **Irrelevance:** Model makes up an answer that doesn't answer the question.
- **Toxicity/Bias:** Model makes up an answer that's harmful/offensive.

# How can LlamaIndex help?

- **Data:** Store additional information beyond raw text chunks.
- **Embeddings:** Optimize embedding representations.
- **Retrieval:** Do better than top-k embedding lookup.
- **Synthesis:** Use LLMs for more than generation.



# How can LlamalIndex help?

---

LlamalIndex provides the following tools:

- **Data connectors** ingest existing data from their native source and format. Could be APIs, PDFs, SQL, and (much) more.
- **Data indexes** structure your data in intermediate representations that are easy and performant for LLMs to consume.
- **Engines** provide natural language access to your data. For example:
  - Query engines are powerful retrieval interfaces for knowledge-augmented output.
  - Chat engines are conversational interfaces for multi-message, “back and forth” interactions with your data.
- **Data agents** are LLM-powered knowledge workers augmented by tools.
- **Application integrations** Tie LlamalIndex into your ecosystem. This could be LangChain, Flask, Docker, ChatGPT, or... anything else!

# Llamaindex with RAG?

When to use Llamaindex with RAG.

## Large Graphs

When working with large graphs representing extensive knowledge bases, Llamaindex can help in efficiently selecting a subset of relevant nodes for retrieval.

This can improve the scalability and speed of the RAG model.

## Diverse Retrieval

Llamaindex is particularly useful when you want diverse retrieval results.

By selecting nodes with diverse content representations, Llamaindex can help ensure that the retrieved information covers a wide range of topics or perspectives.



## Resource Efficiency

Llamaindex allows for more resource-efficient retrieval by focusing on a smaller subset of nodes in the graph. This can be beneficial in scenarios where computational resources are limited or where quick response times are required.

# Llamaindex with RAG?

Llamaindex with RAG might not be appropriate for:

## Specific Retrieval Needs

If your task requires highly specific information that may not be well-represented in the subset selected by Llamaindex, using it might not yield optimal results.

In such cases, other retrieval methods or direct graph traversal techniques may be more suitable.

## Small Graphs

For small graphs where the overhead of applying Llamaindex is unnecessary or where diversity in retrieval results is not a priority, using Llamaindex might be overkill.


Simple retrieval methods or even brute-force approaches could suffice in such scenarios.



## Performance Trade-offs

While Llamaindex can improve efficiency and diversity, it also introduces computational overhead for index construction and maintenance.

If the overhead outweighs the benefits, or if real-time performance is critical, alternative approaches may be preferred.



# Embeddings in Llamalndex

---

# Embeddings in LlamaIndex

---

Embeddings are used in LlamaIndex to represent your documents

- Embedding models take text as input, and return a long list of numbers used to capture the semantics of the text.
- These embedding models have been trained to represent text this way, and help enable many applications, including search!
- At a high level, if a user asks a question about dogs, then the embedding for that question will be highly similar to text that talks about dogs.
- When calculating the similarity between embeddings, there are many methods to use (dot product, cosine similarity, etc.).

# Cohere Chat with Llamaindex

To use Cohere's chat functionality with Llamaindex create a Cohere model object and call the chat function.

```
from llama_index.llms.cohere import Cohere
from llama_index.core.llms.types import ChatMessage

cohere_model = Cohere(api_key="{API_KEY}")

message = ChatMessage(role="user", content= "Who founded Cohere?")

resp = cohere_model.chat([message])

print(resp)
```

# Cohere Embeddings with LlamalIndex

- To use Cohere's embeddings with LlamalIndex create a Cohere Embeddings object with an embedding model from this list and call `get_text_embedding`.

```
from llama_index.embeddings.cohereai import CohereEmbedding

embed_model = CohereEmbedding(
    cohere_api_key="{API_KEY}",
    model_name="embed-english-v3.0", # Supports all Cohere embed models
    input_type="search_query", # Required for v3 models
)

# Generate Embeddings
embeddings = embed_model.get_text_embedding("Welcome to Cohere!")

# Print embeddings
print(len(embeddings))
print(embeddings[:5])
```



# Thank You